

Article

Analysis of Energy Efficiency in Green Cluster Computing

Cathal McStay ^{1,†} and David Cutting ^{2,‡} *

- ¹ Edinburgh Parallel Computing Centre, The University of Edinburgh, Edinburgh, United Kingdom
² School of Electronics, Electrical Engineering and Computer Science, Queen's University Belfast, Belfast, United Kingdom; d.cutting@qub.ac.uk
* Correspondence: d.cutting@qub.ac.uk; Tel.: +44 (0)28 9097 4998
† Both authors conceptualized the work and contributed to the framing and written paper. McStay completed the experimental work, results analysis and initial paper writing under the academic supervision of Cutting.

Abstract: Energy efficiency in computing has emerged as a critical concern due to escalating environmental and financial costs, particularly in the context of cluster computing, where there is an ever-increasing software workload. Achieving meaningful improvements in energy efficiency requires a comprehensive understanding of the interplay between hardware and software. This research investigates how algorithmic optimisations, language choice, and parallelisation strategies influence energy efficiency, and how hardware-level strategies such as underclocking, overclocking, cooling, and on-demand computing further can impact energy usage. Experiments on a custom-built Raspberry Pi Bramble cluster used workloads like Monte Carlo Pi simulations in Python and C. Energy and performance trade-offs were evaluated using Energy Factor (EF), Speed-up, and a new Efficiency-Performance Score (EPS). Results show parallelisation greatly improves energy efficiency over serial execution. Cooling slightly boosts speed under heavy loads but increases total energy use. Perhaps counter-intuitively underclocking actually raises total energy consumption, while overclocking reduces it. Language choice also impacts efficiency, with C offering notable energy savings over Python. The findings support the hypothesis that software optimisation alone can improve energy efficiency, but the most impactful results are achieved when both software and hardware strategies are jointly considered. These insights contribute to the design of future energy-aware computing systems and provide a foundation for sustainable, high-performance computing architectures.

Keywords: Energy Efficiency, Cluster Computing, Green Computing, Raspberry Pi Cluster.

1. Introduction

The rising environmental and financial costs of large-scale computing [1] necessitate energy-efficient system design. This research investigates how hardware and software optimisations affect energy efficiency in cluster computing. The goal is to demonstrate that integrating energy-aware strategies at both levels can enhance performance while reducing energy consumption.

Higher-level languages offer improved usability; however, they often introduce resource overhead compared to lower-level languages [2]. Similarly, the performance gains obtained through parallelisation may improve energy efficiency by reducing execution time, but can also increase total energy consumption if additional power draw is not offset. Additionally, algorithmic improvements in software design can enhance energy efficiency. For instance, reducing the complexity of a program from $O(n^3)$ to a lower order reduces unnecessary computation, thereby conserving energy [3].

From a hardware perspective, strategies such as cooling, clock speed adjustment, and alternative hardware architectures may contribute to energy savings [4]. Cooling components like the CPU and GPU helps prevent thermal throttling, thereby maintaining higher performance and utilisation. However, this can also lead to increased overall energy consumption, as the system operates more efficiently and performs more computations per second, unlike throttled components, which consume less energy as a consequence

Citation: McStay, C. and Cutting, D. Analysis of Energy Efficiency in Green Cluster Computing. *Electronics* **2026**, *1*, 0. <https://doi.org/>

Received:

Revised:

Accepted:

Published:

Copyright: © 2026 by the authors. Submitted to *Electronics* for possible open access publication under the terms and conditions of the Creative Commons Attribution (CC BY) license (<https://creativecommons.org/licenses/by/4.0/>).

of reduced computational activity. Adjusting the clock speed of computing components directly affects the rate at which computations are performed. While increasing the clock speed typically leads to higher power draw, it can also reduce overall energy consumption if the performance gain significantly shortens execution time. Conversely, underclocking may lower power draw but prolong computation, potentially increasing total energy consumed. Additionally, as novel computing paradigms continue to evolve, system architectures may increasingly shift away from the traditional models toward designs that prioritise energy efficiency, such as heterogeneous systems or domain-specific accelerators.

This study bridges the gap between software optimisations and hardware configurations, areas often studied separately. Experiments on a custom Raspberry Pi cluster examine how energy consumption differs depending on the choice of language, the efficiency of the algorithm, parallelisation strategies and implementations, cooling, underclocking, overclocking, and on-demand computing.

The setup consists of three Raspberry Pis running Raspbian Lite, networked via a switch and monitored by an Arduino Uno with an SCT-013 current sensor. A fourth Pi acts as coordinator and energy logger, with workloads (e.g., Monte Carlo Pi) executed using MPI and OMP in Python and C to evaluate energy efficiency.

To calculate continuous energy (the running draw of energy), the following equation can be applied:

$$\text{Watts} = \text{Amps} \times \text{Voltage} \quad (1)$$

To calculate the total energy consumption of a software workload, continuous power readings are taken throughout its execution. Since power consumption can fluctuate based on workload intensity, a weighted average of these readings is computed and then multiplied by the total runtime to estimate energy usage. In this study, workloads varied from a few seconds to several minutes in duration, with some extending to over an hour. As such, energy is expressed in Watt-hours (Wh):

$$\text{Watt-Hour(Wh)} = \text{Watts} \times \text{Hours} \quad (2)$$

To evaluate the impact of different strategies on energy consumption, it is first essential to define what is meant by energy efficiency in this context. Here, energy efficiency refers to the ability to complete a fixed set of computations while using less total energy.

This notion of efficiency can be formalised using the metric Energy Factor (EF), which quantifies the relative change in energy consumption when applying a new strategy compared to a baseline. An EF greater than 1 indicates improved efficiency, whereas an EF less than 1 implies greater energy use for the same output.

$$EF(\text{Energy Factor}) = \frac{\text{Initial Energy Requirements}}{\text{New Energy Requirements}} \quad (3)$$

To complement this, Speedup is used to measure the improvement in execution time:

$$\text{Speedup} = \frac{\text{Previous Time Taken}}{\text{New Time Taken}} \quad (4)$$

Again, a value above 1 indicates a faster execution, while a value below 1 indicates a slowdown. In both metrics, the *desired outcome* is for the new configuration to use *less time or energy*, resulting in a score greater than 1.

A simple combined metric could be a direct division of EF and Speedup:

$$EPS_{\text{simple}} = \frac{EF}{Speedup} \quad (5)$$

However, this approach may be overly influenced by extreme values in either energy or performance. To address this, a geometric mean is used, resulting in a more balanced metric called the Efficiency-Performance Score (EPS):

$$EPS = \sqrt{EF \times \text{Speedup}} \quad (6)$$

This formulation ensures that both energy efficiency and performance contribute equally to the final score, preventing one from disproportionately dominating the result.

For cases where energy efficiency is prioritised over performance (e.g., in sustainability-focused applications), a Weighted Efficiency-Performance Score (WEPS) can be used:

$$WEPS = \sqrt{EF \times (\text{Speedup} \times w)} \quad (7)$$

Where $w \in (0, 1)$ is a weighting factor that reduces the influence of speedup. For example, setting $w = 0.5$ places greater emphasis on energy savings.

It is anticipated that results will highlight the interdependency of software and hardware optimisations, revealing that neither approach alone can fully address energy inefficiency. Instead, a combined strategy is expected to yield significant reductions in energy consumption without compromising performance. Moreover, applying both software optimisations (e.g. parallelisation and algorithm refinement) and hardware strategies (e.g. overclocking and cooling) may reveal a positive correlation between energy efficiency and performance.

2. Literature Review

This literature review covers key aspects of green computing, including its principles and importance. It examines the environmental and financial impacts of computing, highlights hardware and infrastructural strategies to reduce energy consumption, and explores software-level optimisations for energy efficiency in cluster computing. In addition to this, few studies directly compare software and hardware optimisations working in tandem, there are gaps in the existing literature and this will be made clear, as well as proposed approaches for addressing these gaps.

2.1. Green Computing

Green computing focuses on designing and optimising computer systems, networks, and software to maximise energy efficiency and reduce environmental impact [5]. Notably, none of the top 10 systems on the recent GREEN500 list appear in the TOP500 list's top 10 [6].

Why does green computing matter? Environmental impacts arise throughout the computing lifecycle, from resource extraction and hardware manufacturing to energy consumption during operation. The HPC sector alone consumes approximately 5.2 TWh of energy annually, equivalent to a 600MW load [4]. According to a 2023 UK government report, the carbon intensity of electricity was 0.225 kgCO₂e/kWh [7], meaning that potentially, HPC systems contribute around 1.17 million metric tonnes of CO₂ yearly about 0.003% of the global total of 37.4 billion tonnes [8]. However, this figure does not take into account the number of organisations which use sustainable methods for powering their HPC centres, for example, Google, Facebook and Microsoft claim to have renewable energy consumption at 100%, 50% and 50% respectively [9]. In addition to the companies above, Amazon also claim to use 100% renewable energy [10]. Meanwhile, the trend of countries as a whole moving towards renewables continues, 2019, China, Germany, and the USA used 22%, 40% and 17% renewable energy [11]. In the paper by Strubell et al. they claim that a more accurate way of measuring CO₂ would be taking this into account resulting in the measurement, CO₂ = 0.477 pounds per kWh which would result in a figure closer to 0.22 million metric tonnes yearly. However, this estimation is based on publicly available data from companies which may obfuscate the actual figures. There is a gap in research into more accurate measurements for the amount of carbon emissions from electricity used in computing which is understandable due to the dynamic nature of energy production. This lack of transparency complicates accurate assessment of carbon emissions associated with computing workloads.

The broader computing industry also has a significant environmental footprint. In 2022, global data centres consumed an estimated 240-340 TWh of electricity, representing about 1-1.3% of global final electricity demand [1]. This figure excludes cryptocurrency mining, which added approximately 110 TWh, accounting for 0.4% of global electricity consumption. Combined, there is a potential 101.25 million metric tonnes of CO₂ being produced, representing nearly 0.3% of global CO₂ production. Despite efficiency improvements, rising workloads in large data centres have driven a 20-40% annual increase in energy use in recent years [1]. There will also be a further rise in the energy usage of large computation centres, with the evolution of LLMs, the energy used to train these models and then used to allow for requests to be sent to these models is increasing. The models introduced by OpenAI began with 117 million parameters in 2018, while the latest model has more than a trillion parameters, the previous model GPT-3, was estimated to require 1287 MWh, and in January 2023 the energy consumption of ChatGPT was projected to be over 1500 MWh [12]. These energy costs will only rise with the training of new models and necessity for more hardware, specifically GPUs. In addition to this, more organisations are looking at developing their own models for their specific use cases, which will require specialised training per model. The costs for training LLMs and then the further costs of hosting and exploiting these new technologies is becoming difficult to find, as industry attempts to shore up competitive moats and restrict information regarding their underlying LLM technologies, these details can become less reliable and available [13].

Beyond environmental concerns, financial costs are significant. Rising energy prices increase the need for either more funding or improved efficiency. It is difficult to obtain an average statistic for the cost of energy worldwide, an average may be calculated by examining the statistics provided by different organisations. The EU provides an average non-household consumer cost of \$0.21 per kWh for the first half of 2024 [14]. The US Bureau of Labour provides an average of \$0.179 [15]. Meanwhile, in China the average cost for businesses is \$0.87 [16]. This results in an average of \$0.42 per kWh. While this method of calculating the energy costs of computing is not ideal, it provides a useful statistic. The annual energy expense for HPC systems totals approximately £1.664 billion, while the wider community, including the costs for crypto-mining could reach £144 billion.

2.2. Physical strategies for reducing energy consumption

There are many physical and hardware strategies for reducing energy consumption. Using liquid cooling systems to reduce heat and save on cooling costs is one measure. For example, Microsoft's experiment with immersion cooling at its data centres resulted in a 90% reduction in cooling energy consumption illustrating the potential of this technology. Furthermore, the transition to green refrigerants, like ammonia or CO₂, highlights a conscious effort to reduce the ecological impact of cooling systems [4].

Another promising approach is hardware-enforced power bounding, as discussed in "Beyond DVFS: A First Look at Performance Under a Hardware-Enforced Power Bound" [17]. This technique involves dynamically limiting processor power consumption through mechanisms such as Intel's Running Average Power Limit (RAPL). By setting power limits, the hardware adjusts CPU performance to remain within the specified power budget, enabling predictable energy use. However, this can cause performance variation due to differences in processor efficiency. Additionally, power clamping emerges as a powerful alternative to traditional DVFS, offering more precise energy management by directly limiting the power a processor can draw during operation. This approach can be dynamically adjusted based on workload demands, thus ensuring system stability while controlling energy consumption [17].

In FPGA-based systems, combining undervolting and overclocking can improve energy efficiency by reducing static and dynamic power consumption while enhancing performance, as shown in [19]. However, this approach risks bit-level faults in critical components like on-chip memory and data paths, requiring robust error correction and hardware design. Moreover, this is a specific use case when most HPC clusters do not




Figure 1. The CPU is a major power consumer in a system. Techniques such as DVFS can help reduce the energy usage of both the CPU and memory [18].

feature FPGA devices and the paper [19] doesn't consider the impact of following the above on traditional computing systems.

Fine-grained tuning of supply voltages and clock speeds can cut power usage by up to 60% while keeping error rates manageable. Adaptive power management further optimises performance in real-time applications. Similar energy-saving techniques have also been applied to modern HPC processors, achieving significant gains with minimal tuning effort [20].

Another strategy for saving energy is to suspend or power down idle components or nodes using techniques such as sleep modes, Wake-on-LAN (WoL), or hibernation. As noted by [21], “Since the WoL technology is common in modern PCs, employing EnergySave in networks of PCs can represent a cheap solution that can bring high savings under the energy (kWh), economic (€), and greenhouse gases (tons of CO₂).” More advanced approaches include Intel’s Demand-Based Switching, which dynamically scales down CPU clock speeds or powers off idle chip regions to reduce power consumption [22].

Incorporating heterogeneous architectures, such as GPUs and other accelerators alongside traditional CPUs, has been shown to enhance energy efficiency in HPC environments. GPUs, designed for parallel processing, offer significant energy savings over CPUs, particularly in applications that benefit from massive parallelism. Studies have demonstrated that for simple computational tasks, GPUs can reduce energy consumption by a factor of 1.1 to 3.2 times compared to CPUs [23]. As the complexity of tasks increases, FPGAs outperform both CPUs and GPUs, achieving energy reduction ratios of 1.2 to 22.3 times for complex vision pipelines [23]. However, these figures come from a study which focuses on the specific use-case of computer vision.

In contrast, CPU-only solutions, while versatile, often consume more energy due to their sequential processing nature, which is often less energy-efficient for workloads that benefit from parallel execution. A holistic approach, like the GreenGPU framework, has been proposed to optimize energy efficiency by dynamically splitting workloads between GPUs and CPUs, and by adjusting the frequencies of both components based on their utilization. This method has been shown to achieve up to 21.04% energy savings [24]. By utilising the strengths of each hardware type, GPUs for high-throughput parallel tasks and CPUs for more sequential tasks, heterogeneous architectures can significantly reduce energy consumption without compromising performance. Therefore, integrating GPUs and other accelerators into HPC systems represents a promising direction for achieving green computing goals while maintaining computational power.

2.3. Software strategies for reducing energy consumption

Software optimisation, coupled with hardware enhancements, is essential. Effective algorithms, particularly in fields such as deep learning, demonstrate the ability to diminish energy usage considerably. Code optimisation can prompt reductions in computation time and thus lower energy consumption. Studies have indicated that software-level enhancements may reduce energy consumption by up to 50%, highlighting the latent

potential in this area [4]. However, it is important to note that these indications focus on their own software and will vary depending on the software and system in question.

The simple rule that Energy = Power × Time, suggests two general ways to save energy: 1) faster speed, given a constant power; and 2) lower power without increasing run time [25]. This would lead to the belief that in all situations decreasing the time through increased performance. However, improving performance and efficiently managing power may conflict with each other because faster speeds frequently come from using more resources less efficiently, which may excessively increase power consumption [25].

One software strategy for reducing energy consumption is using less accurate data types which incur less resource cost where the accuracy of data can be sacrificed slightly for improved energy efficiency [26]. Through porting applications such as Monte Carlo Pi to use approximations of nearly all floating point operations, there were gains of up to 15% in energy efficiency. However, the paper [26] failed to give an example of software for which reducing the accuracy of data may have a larger impact on the *correctness* of the output, so this should be a consideration in any development plan.

Virtualisation is an effective strategy for partitioning hardware resources, enabling multiple services to run within virtual machines (VMs) on a single physical system. This consolidation increases overall energy efficiency by reducing the need for additional hardware. As multiple VMs share the same hardware, utilisation improves, and energy demands decrease, particularly for cooling, since fewer physical machines are required to deliver the same functionality [22].

At the organisational level, scheduling plays a crucial role in improving the energy efficiency of data centres. Liu et al. explored this by designing an energy-aware scheduler for data grid systems that support both real-time and data-intensive applications. Their approach considers both the physical location of data and the characteristics of the applications to optimise scheduling decisions. The proposed distributed scheduler integrates task scheduling with data placement strategies to minimise energy usage. Significant energy savings are achieved by reducing unnecessary data replication and limiting task migration across nodes [22,27].

2.4. Gaps in the existing literature and how this approach differs

Despite the depth of research into hardware and software energy-saving techniques, a key gap lies in the lack of empirical studies that examine both in tandem. Most literature treats hardware and software strategies independently, leaving open questions about their combined effects or synergies. Additionally, industry confidentiality often limits access to accurate data on real-world energy consumption and optimisation effectiveness, further complicating cross-layer comparisons. However, it should be considered that this lack of an overarching review, may be due to the uniqueness of energy efficiency in individual projects; a problem which is embarrassingly parallelisable may lead to more energy savings than one which requires more sequential operations, for example. In contrast, this review lays the groundwork for an experimental evaluation of cross-layer strategies using a Raspberry Pi cluster, a novel test bed that facilitates real-time measurement and control of both software and hardware parameters.

3. Technical Approach

This section outlines the methodology used to investigate energy efficiency within a small-scale cluster computing environment. The goal is to quantify how software and hardware optimisations affect energy usage and performance using controlled, repeatable experiments.

3.1. Investigating the Problem

As high-performance computing continues to scale, energy consumption has emerged as a critical bottleneck, not just in terms of cost, but also in terms of environmental impact. Measuring energy efficiency is no longer optional; it is a necessary step toward sustainable

system design. This project seeks to investigate how and why different decisions in software and hardware affect energy usage, and whether gains in performance come at the cost of higher power draw. For the purposes of this work, energy efficiency in computing is defined as the change in total energy consumption for completing the same computational workload.

While Raspberry Pis are not directly representative of industrial HPC clusters, they allow cost-effective, scalable experimentation, particularly suitable for investigating trends related to parallelisation, energy consumption, and thermal effects. Limitations such as Amdahl's Law and thermal throttling are considered when interpreting results.

3.2. Assumptions and Experimental Constraints

The voltage supply to the cluster is assumed to be stable at 230V throughout testing. All tests are conducted with identical operating system images; system load at boot time is minimised, and no background services are allowed to run during workloads. Within the work, we distinguish between *compute/node energy*, which is measured, and total *system-level energy*, for example including cooling, which we do not.

The Arduino-based current sensor setup introduces minor latency between sampling and logging, which may slightly affect precision for extremely short workloads. However, due to the averaging of current over longer intervals, the overall effect is expected to be negligible.

In order to ensure a fair and consistent measurement of energy usage, baseline power values were established for the Raspberry Pi cluster at various operational stages. At idle, each Raspberry Pi was observed to draw approximately $3.27 \text{ W} \pm 10\%$, resulting in a combined idle consumption of around $9.81 \text{ W} \pm 10\%$ for the three-node cluster. These values were calculated using current readings obtained from the SCT-013 current sensor and Arduino setup and confirmed by the Tapo smart plug.

In contrast, when all four CPU cores were fully utilised on each of the three Raspberry Pis (using CPU stress tests), the total power draw increased significantly, ranging between 16 and 18 W. This value represents the system under sustained computational load and serves as the upper bound for evaluating the energy efficiency of software workloads under test conditions.

3.3. Measuring Energy Usage

Quantifying energy usage is essential for evaluating the impact of different software and hardware strategies. In this project, energy consumption is calculated based on electrical power, derived from the product of current and voltage, as shown in Equation 1. The power metric is then combined with execution time (Equation 2) to determine total energy consumed. For the scale of experiments conducted on the Raspberry Pi cluster, energy usage is most commonly reported in watt-hours.

To measure current in real time, an SCT-013 non-invasive current sensor is used in conjunction with an Arduino Uno. The sensor clamps onto the live wire of the Raspberry Pi cluster's power strip and outputs analogue current readings to the Arduino, which are then streamed via serial connection to a coordinating Raspberry Pi node. The accuracy of the SCT-013 sensor has been validated by comparing its readings with those from a Tapo smart plug and validated against the expected demand from the hardware specification. The close alignment of values across idle and workload states supports the reliability of the chosen measurement setup.

This measurement pipeline is tightly integrated with a custom job-runner script that automates workload execution and synchronises energy logging. When a job is submitted, the script initiates the energy monitoring process, executes the selected workload (e.g., Monte Carlo Pi), records runtime, and logs corresponding power readings. These logs are then parsed to compute average power, total energy consumed, and efficiency metrics such as Energy Factor (EF) and Efficiency-Performance Score (EPS). This integration enables

consistent, low-overhead experimentation and ensures accurate energy tracking across test scenarios.

3.4. Hardware

The experimental setup consists of a Raspberry Pi “bramble” cluster composed of three compute nodes and one coordinator node. Each Raspberry Pi is connected via Ethernet through a central switch to facilitate low-latency communication and synchronised workload distribution. The coordinator Pi also handles energy measurement and logging from the Arduino-based current sensor system. A visual schematic of the cluster setup is shown in Figure 2.




Figure 2. Raspberry Pi Bramble Cluster Topology

To investigate thermal effects on performance and energy usage, an external cooling fan is used during selected runs. While the Raspberry Pi does not have active cooling by default, under high computational load, the CPU is prone to thermal throttling, which reduces clock frequency and, in turn, performance. Cooling mitigates this throttling, potentially improving runtime but at the cost of additional energy as the nodes are able to perform more computation.

Clock speed is another key hardware variable explored in this study. Underclocking reduces the maximum frequency of the CPU, which theoretically lowers power draw but may lead to longer execution times that negate energy savings. Conversely, overclocking increases performance, potentially reducing runtime and total energy used, provided the increase in power is offset by the decrease in time. However, overclocking may also increase thermal output, leading to throttling or instability if not properly managed. Moreover, as will be encountered later, there are also stability issues when changing the clock speed of components. The Raspberry Pi 3B+ which was employed in this setup has a recommended clock speed of 1.4 GHz and through testing, operating at 1.6 GHz resulted in system instability on multiple nodes. Therefore, the minimum clock speed of 0.6 GHz to 1.5 GHz will be tested.

Another strategy tested is “on-demand computing”, where idle nodes are powered off when not participating in a computation. In principle, this reduces idle energy draw, but in practice, it introduces challenges in coordination, system responsiveness, and reinitialisation time. Additionally, the energy cost of restarting or reconfiguring nodes may outweigh savings from temporary shutdowns in short workloads.

Although Raspberry Pis are not directly comparable to industrial HPC hardware in terms of performance, memory, or thermal management, they provide a scalable and accessible platform for prototyping and controlled experimentation. The insights gained, while not directly transferable in magnitude, are valuable in understanding trends in energy-performance trade-offs and can inform energy-aware practices in larger, more powerful cluster systems.

3.5. Software

Each Raspberry Pi in the cluster runs Raspbian OS Lite, a lightweight Debian-based distribution optimised for performance and minimal overhead. This ensures that background processes do not distort energy measurements. Communication between nodes is handled via Secure Shell (SSH), while distributed workloads are executed using the Message Passing Interface (MPI). For shared-memory parallelism within individual nodes, OpenMP is employed.

The role of software design in energy efficiency is central to this research. Optimisation can occur at several levels — algorithmic (e.g. reducing computational complexity), architectural (e.g. multithreading and workload distribution), and linguistic (e.g. using low-level languages with minimal runtime overhead). To explore these dimensions, all workloads are implemented in both Python and C. While Python offers ease of development and high readability, it introduces significant overhead due to its interpreted nature. In contrast, C provides lower-level memory and execution control, often resulting in faster runtimes and reduced energy consumption.

The Monte Carlo Pi algorithm estimates π by generating random points within a square and counting how many fall inside a quarter circle. It is used as the primary benchmark for parallel performance due to its embarrassingly parallel nature. The algorithm scales well across cores and nodes, requiring minimal communication overhead.

To complement this, the recursive Fibonacci algorithm, which computes each term as the sum of the two preceding terms, is used to represent more sequential workloads. Comparing it with Monte Carlo Pi enables an evaluation of how workload parallelisation influences energy efficiency under different configurations.

Graph traversal algorithms including Breadth-First Search (BFS), which explores nodes level by level; Depth-First Search (DFS), which explores as deep as possible before backtracking; and Dijkstra's algorithm, which finds the shortest paths from a source node to all others in a weighted graph are also included. These reflect real-world data centre workloads such as web crawling, social network analysis, and routing. Their performance is assessed in both sequential and parallel contexts.

Finally, sorting algorithms such as Bubble Sort, Merge Sort, and Quick Sort are applied to identical randomised arrays. These allow for direct comparisons of energy consumption across different computational complexities. By observing the relationship between input size, algorithmic efficiency (Big-O notation), and actual energy usage, the study reveals how theoretical complexity translates into real-world power consumption.

The variation in algorithmic complexity is visualised in Figure 3

4. Results and Analysis

4.1. Results

4.1.1. Impact of Language

Language choice plays a significant role in the energy efficiency of software. In this project, the same Monte Carlo Pi approximation algorithm was implemented in both Python and C to assess the effect of language-level optimisation on execution time and energy usage. Python, while easy to develop and maintain, is an interpreted language and introduces considerable overhead due to its runtime environment. C, on the other hand, is compiled to native machine code and provides lower-level memory and execution control, resulting in faster runtimes and reduced energy consumption.





Figure 3. Time complexity growth for selected algorithms

Table 1. Monte Carlo Pi – Python Results

N	Time (s)	Energy (Wh)
10M	46.60	0.0062
100M	469.61	0.0631
200M	927.92	0.1220
300M	1392.86	0.2126
400M	1892.58	0.2695
500M	2320.19	0.2553
750M	3482.23	0.5524
900M	4171.48	0.6441

Table 2. Monte Carlo Pi – C Results

N	Time (s)	Energy (Wh)
10M	2.41	0.0003
100M	24.01	0.0034
200M	48.03	0.0058
300M	72.04	0.0104
400M	96.08	0.0140
500M	120.06	0.0134
750M	178.22	0.0232
900M	216.19	0.0363
1B	240.15	0.0309

**Figure 4.** Execution Time Comparison of Python vs C for Monte Carlo Pi




Figure 5. Energy Usage Comparison of Python vs C for Monte Carlo Pi

This result highlights that even for compute-light workloads, compiled languages like C can provide significant energy savings. Over repeated runs or at scale' the efficiency gap between high-level and low-level languages becomes environmentally and financially significant.

4.1.2. Impact of parallelisation

To investigate the role of parallelisation in energy efficiency, two distinct workloads were compared: a non-parallelisable recursive Fibonacci sequence and an embarrassingly parallel Monte Carlo Pi approximation. As both scripts offer differences in computational length, their input values were tuned such that the execution time on a single core was approximately the same, providing a fair baseline for comparison. Specifically, fib(10000000) took 34.01 seconds, while a Monte Carlo Pi approximation with 175 million points took 35.69 seconds on a single core. Once a baseline was established, both workloads were executed across 1 to 12 cores, measuring total time, energy consumed, and derived metrics such as EF, EPS, and Parallel Efficiency. The results in Table 4 clearly show that Monte Carlo Pi benefits significantly from parallelisation. At 12 cores, it achieved a 5.43 \times speedup and a peak EPS of 1.71, indicating that parallel execution not only reduced runtime, but also led to a higher energy saving. In contrast, Fibonacci as shown in Table 3 diminishing returns as cores increased, speedup actually decreased slightly, and energy usage rose. By 12 cores, Fibonacci's EPS dropped to just 0.23, compared to Monte Carlo's 1.57 at the same point. This highlights a key insight: workloads with minimal inter-process communication and independence (like Monte Carlo, where placing dots on a square does not require previous dots or where the other dots are to calculate Pi) are much more energy-scalable across multiple cores. Conversely, inherently sequential workloads (like Fibonacci recursion) suffer from overhead, resource contention, and idle cycles during parallel execution, making them less suitable for energy-efficient parallelism. It is also observed that while the general trend in the Monte Carlo workload, the EPS increases with increasing cores, that due to the smaller workload of 175 million (Monte Carlo Pi can go for much larger values) there is not only a flattening out as the workload approaches 12 cores in terms of EPS, it can be seen there is a decrease for 12 cores in Figure 8

Table 3. Fibonacci for 1,000,000: Energy, Performance and Efficiency Metrics

Processors	Time (s)	Energy (Wh)	Speedup	Parallel Efficiency	EPS
1	34.01	0.067083	1.000	1.000	1.000
2	35.07	0.072003	0.970	0.485	0.686
3	37.06	0.074239	0.918	0.306	0.530
4	38.53	0.079158	0.883	0.221	0.441
5	40.94	0.085419	0.830	0.167	0.372
6	41.15	0.087208	0.827	0.138	0.337
7	41.61	0.106886	0.818	0.117	0.309
8	41.32	0.104203	0.823	0.103	0.291
9	41.34	0.098389	0.823	0.091	0.274
10	41.25	0.100625	0.824	0.082	0.261
11	41.36	0.100625	0.822	0.075	0.248
12	41.84	0.106439	0.813	0.068	0.235

Table 4. Monte Carlo Pi for 175,000,000 Dots: Energy, Performance and Efficiency Metrics

Processors	Time (s)	Energy (Wh)	Speedup	Parallel Efficiency	EPS
1	35.69	0.070661	1.000	1.000	1.000
2	18.36	0.037119	1.943	0.971	1.375
3	13.23	0.025044	2.698	0.900	1.557
4	10.55	0.021467	3.383	0.845	1.695
5	10.51	0.022361	3.395	0.680	1.586
6	9.16	0.020572	3.896	0.650	1.591
7	8.12	0.021019	4.395	0.628	1.662
8	7.44	0.0161	4.798	0.600	1.696
9	6.99	0.014311	5.106	0.568	1.701
10	6.57	0.0161	5.432	0.543	1.718
11	6.39	0.020125	5.586	0.508	1.684
12	6.57	0.021467	5.432	0.453	1.568




Figure 6. Speedup of Fibonacci and Monte Carlo Pi across increasing core counts




Figure 7. EF across cores for Fibonacci and Monte Carlo Pi




Figure 8. EPS for Fibonacci and Monte Carlo Pi across core counts

4.1.3. Impact of Cooling

Cooling reduces thermal throttling and should, in theory, improve sustained performance. However, any performance gain must be weighed against the fact that cooling systems themselves consume additional energy, a factor not measured in this experiment and is therefore excluded from the energy analysis but acknowledged as a factor.

In practice, improved cooling may lead to two outcomes: either (1) execution time decreases, and thus overall energy consumption drops, or (2) performance increases at the cost of higher power draw, resulting in greater total energy usage. This section explores this trade-off using two Monte Carlo workloads (with $N = 10^9$ and $N = 10^{10}$) executed, comparing cooled vs. non-cooled conditions across core counts from 1 to 12. The motivation for testing this configuration with input = 10^9 and 10^{10} was to increase the time spent in operation, therefore leading to an increased likelihood of thermal throttling, the ability to check whether or not a node is throttling comes from the CLI of that node which gives the temperature and whether the node is throttling.

Table 5. Monte Carlo Pi ($N = 1,000,000,000$): Non-Cooled System

Processors	Time (s)	Energy (Wh)	Speedup	Parallel Eff.	EF	EPS
1	452.37	1.091669	1.000	1.000	1.000	1.000
2	227.78	0.574233	1.986	0.993	1.901	1.943
3	152.67	0.399369	2.963	0.988	2.733	2.846
4	115.97	0.317975	3.901	0.975	3.433	3.659
5	93.23	0.276831	4.852	0.970	3.943	4.374
6	78.50	0.232556	5.763	0.960	4.694	5.201
7	67.07	0.188728	6.745	0.964	5.784	6.246
8	59.25	0.169050	7.635	0.954	6.458	7.022
9	53.01	0.151161	8.534	0.948	7.222	7.850
10	48.05	0.150714	9.414	0.941	7.243	8.258
11	44.09	0.139086	10.260	0.933	7.849	8.974
12	40.06	0.131483	11.292	0.941	8.303	9.683

Table 6. Monte Carlo Pi ($N = 1,000,000,000$): Cooled System

Processors	Time (s)	Energy (Wh)	Speedup	Parallel Eff.	EF	EPS
1	454.08	1.207053	1.000	1.000	1.000	1.000
2	227.88	0.519225	1.993	0.996	2.325	2.152
3	152.87	0.399369	2.970	0.990	3.022	2.996
4	118.82	0.311267	3.822	0.955	3.878	3.850
5	94.04	0.266097	4.829	0.966	4.536	4.680
6	78.37	0.232108	5.794	0.966	5.200	5.489
7	67.46	0.199014	6.731	0.962	6.065	6.389
8	59.33	0.169944	7.653	0.957	7.103	7.373
9	52.89	0.159211	8.585	0.954	7.581	8.068
10	47.38	0.144900	9.584	0.958	8.330	8.935
11	43.58	0.144006	10.419	0.947	8.382	9.345
12	39.82	0.129247	11.403	0.950	9.339	10.320

Table 7. Monte Carlo Pi ($N = 10,000,000,000$): Cooled System

Processors	Time (s)	Energy (Wh)	Speedup	Parallel Eff.	EF	EPS
1	4523.46	13.390281	1.000	1.000	1.000	1.000
2	2263.03	7.438647	1.999	0.999	1.800	1.897
3	1509.00	5.143503	2.998	0.999	2.603	2.794
4	1135.02	3.936897	3.985	0.996	3.401	3.682
5	909.89	3.169017	4.971	0.994	4.225	4.583
6	757.12	2.464642	5.975	0.996	5.433	5.697
7	650.15	2.252658	6.958	0.994	5.944	6.431
8	568.36	1.834953	7.959	0.995	7.297	7.621
9	507.90	1.680661	8.906	0.990	7.967	8.424
10	456.60	1.682897	9.907	0.991	7.957	8.878
11	415.35	1.510717	10.891	0.990	8.864	9.825
12	379.40	1.412775	11.923	0.994	9.478	10.630

Table 8. Monte Carlo Pi ($N = 10,000,000,000$): Uncooled System

Processors	Time (s)	Energy (Wh)	Speedup	Parallel Eff.	EF	EPS
1	4525.15	11.431	1.000	1.000	1.000	1.000
2	2263.54	5.787056	1.999	1.000	1.975	1.987
3	1513.94	4.00085	2.989	0.996	2.857	2.922
4	1157.78	3.659619	3.908	0.977	3.124	3.494
5	922.56	2.802294	4.905	0.981	4.079	4.473
6	766.29	2.482531	5.905	0.984	4.605	5.215
7	657.73	2.043806	6.880	0.983	5.593	6.203
8	583.01	1.939603	7.762	0.970	5.893	6.763
9	518.74	1.697656	8.723	0.969	6.733	7.664
10	467.26	1.550072	9.684	0.968	7.374	8.451
11	424.41	1.390861	10.662	0.969	8.219	9.361
12	389.27	1.400253	11.625	0.969	8.164	9.742

Across the cooled and uncooled workloads for $N = 10^9$ it is evident that there is a slight increase in performance for growing processor count when the bramble is being cooled. This suggests that while a small number of processors are pegged this does not massively increase the temperature on the chip, however, when all processors are pegged this increases the temperature to the point where they are throttling and, therefore, gain more of a speedup. For $N = 10^{10}$, this difference is more pronounced as with increasing

amounts of time spent in the workload there is an increase in the temperature and this causes throttling and as such cooling the system will result in a better speedup. In terms of EF, as can be viewed in Figure 9, there is a clear trend that for rising values of processors the EF increases showing an increased energy efficiency for an increasing number of processors, it is clear that the gains found in the decrease of time spent working outweighs the higher energy use through calculations. Moreover, for EPS in Figure 10, this trend is similar to EF, however, as this stat also takes speedup into account, it can also prove that there is a positive correlation in this case between performance and energy efficiency.

451
452
453
454
455
456
457
458




Figure 9. EF Comparison across Cooled and Uncooled Monte Carlo Pi Executions ($N = 10^9$ and $N = 10^{10}$)




Figure 10. EPS Comparison across Cooled and Uncooled Monte Carlo Pi Executions ($N = 10^9$ and $N = 10^{10}$)

4.1.4. Impact of On-Demand Computing

On-demand computing refers to dynamically powering nodes only when required, rather than keeping all resources powered continuously. While there is some overhead in spinning resources up or down, for predictable workloads with known resource requirements, this model can significantly reduce energy waste from idle components. In terms of pure energy usage as can be seen in Figure 11, the actual amount of energy used for not powering all 3 nodes, provides a much higher energy savings. However, due to the way that the measurements created for this research operate, $\frac{\text{energy on 1 core}}{\text{energy on } n \text{ cores}}$ and EPS which multiplies speedup and EF, due to the much lower starting energy values (due to only powering one node) the EF and EPS for on-demand are much lower than that of non-demand, even though the trend in speedup is the same. This experiment does show, however, that for workloads which do not require all of the nodes, on-demand computing is a potential solution for energy efficiency.

Table 9. Monte Carlo Pi (N = 1,000,000,000): On-Demand System

Cores	Time (s)	Energy (Wh)	Speedup	Parallel Eff.	EF	EPS
1	453.76	0.321106	1.000	1.000	1.000	1.000
2	227.81	0.175758	1.9918	0.9959	1.8270	1.9076
3	152.52	0.125222	2.9751	0.9917	2.5643	2.7621
4	116.24	0.101967	3.9036	0.9759	3.1492	3.5061
5	93.57	0.162789	4.8494	0.9699	1.9725	3.0928
6	77.83	0.141769	5.8301	0.9710	2.4165	4.6339
7	66.80	0.125222	6.7921	0.9704	2.5643	4.1736
8	58.73	0.108576	7.7262	0.9658	2.9544	4.7779
9	53.08	0.174864	8.5486	0.9498	1.8363	3.9621
10	47.19	0.156528	9.6156	0.9616	2.0514	4.4414
11	43.20	0.148031	10.5037	0.9549	2.1692	4.7733
12	39.83	0.134614	11.3924	0.9494	2.3854	5.2129

Table 10. Monte Carlo Pi (N = 1,000,000,000): Non-Demand System

Cores	Time (s)	Energy (Wh)	Speedup	Parallel Eff.	EF	EPS
1	454.08	1.207053	1.000	1.000	1.000	1.000
2	227.88	0.519225	1.9926	0.9963	2.3247	2.1523
3	152.87	0.399369	2.9704	0.9901	3.0224	2.9963
4	118.82	0.311267	3.8216	0.9554	3.8779	3.8496
5	94.04	0.266097	4.8286	0.9657	4.5361	4.6801
6	78.37	0.232108	5.7941	0.9657	5.2004	5.4892
7	67.46	0.199014	6.7311	0.9616	6.0652	6.3895
8	59.33	0.169944	7.6535	0.9567	7.1027	7.3729
9	52.89	0.159211	8.5854	0.9539	7.5815	8.0678
10	47.38	0.144900	9.5838	0.9584	8.3302	8.9351
11	43.58	0.144006	10.4195	0.9472	8.3820	9.3453
12	39.82	0.129247	11.4033	0.9503	9.3391	10.3197




Figure 11. Energy Usage Comparison between On-Demand and Non-Demand systems




Figure 12. EF Comparison between On-Demand and Non-Demand systems





Figure 13. EPS Comparison between On-Demand and Non-Demand systems

4.1.5. Impact of Clocking

Adjusting the clock speed of processors, either through underclocking or overclocking, offers a direct mechanism for influencing both performance and energy consumption. At higher clock speeds, processors can execute instructions more rapidly, reducing the overall execution time. However, this typically comes at the cost of increased power draw per unit time. Conversely, underclocking reduces energy consumption per second but prolongs the execution time of workloads. To investigate this trade-off, a Monte Carlo Pi workload ($N = 10^9$) was executed on all 12 cores of a Raspberry Pi 3B+ cluster under different clock frequencies, ranging from 600 MHz to 1500 MHz. As shown in Table 11, increasing the clock speed consistently improved both speedup and energy efficiency metrics (EF and EPS), indicating that the gains in performance outweighed the rise in energy consumption per second. Notably, the data reveals that the fastest configuration (1500 MHz) resulted in the best EPS value (4.08), suggesting that the shorter execution time more than compensated for the increased instantaneous power draw. The EF metric also increased with clock speed, demonstrating improved energy-to-performance scaling. However, attempts to increase the clock speed beyond 1500 MHz caused instability—devices failed to boot and required reflashing. This limitation highlights a practical ceiling for overclocking on the Raspberry Pi 3B+ hardware used in this study.

Table 11. Impact of Clock Speed on Energy and Performance (12 Cores, $N = 10^9$)

Clock Speed (MHz)	Time (s)	Energy (Wh)	Speedup	EF	EPS
600	40.19	0.118961	1.000	1.000	1.000
700	34.62	0.108228	1.1609	1.0992	1.2760
800	30.63	0.097942	1.3121	1.2146	1.5937
900	27.87	0.085867	1.4421	1.3854	1.9978
1000	24.51	0.084972	1.6397	1.4000	2.2956
1100	22.54	0.079606	1.7831	1.4944	2.6645
1200	20.77	0.074239	1.9350	1.6024	3.1007
1300	19.93	0.076028	2.0166	1.5647	3.1553
1400	18.24	0.075133	2.2034	1.5833	3.4887
1500	17.12	0.068425	2.3475	1.7386	4.0814

**Figure 14.** Impact of Clock Speed on Energy Consumption, EF, and EPS

4.1.6. Sorting Algorithms and the impact of Big O Complexity

Sorting algorithms were selected to empirically explore how algorithmic complexity impacts both execution time and energy usage. Sorting algorithms provide well-defined theoretical complexities:

- Bubble Sort: $\mathcal{O}(n^2)$ simple but inefficient.
- Merge Sort: $\mathcal{O}(n \log n)$ efficient.
- Quick Sort: Quick Sort: $\mathcal{O}(n \log n)$ on average, fast but can degrade to $\mathcal{O}(n^2)$.

This made them ideal candidates to test whether energy profiles align with theoretical performance expectations. The results below show that Bubble Sort performs significantly worse, with its execution time and energy usage ballooning as input size N increases. Merge and Quick Sort remain efficient even at larger N , both in terms of runtime and Wh consumed.

Table 12. Bubble Sort Results

N	Time (s)	Energy (Wh)
10	2.52	0.000106
100	1.23	0.000036
1000	1.22	0.000036
10000	3.45	0.000115
50000	221.27	0.008403
100000	5531.56	0.220071

Table 13. Merge Sort Results

N	Time (s)	Energy (Wh)
10	1.15	0.000036
100	1.15	0.000036
1000	2.41	0.000080
10000	1.19	0.000036
50000	1.46	0.000036
100000	2.82	0.000115

Table 14. Quick Sort Results

N	Time (s)	Energy (Wh)
10	1.19	0.000036
100	1.15	0.000036
1000	1.25	0.000036
10000	1.22	0.000036
50000	1.54	0.000080
100000	2.60	0.000115




Figure 15. Execution Time for Sorting Algorithms




Figure 16. Energy Usage for Sorting Algorithms

4.1.7. Graph Algorithms and the impact of traversing increasing number of nodes

Graph traversal algorithms are critical in many domains of HPC, from network analysis to optimisation problems. This experiment evaluated three classic algorithms, Dijkstra, Breadth-First Search (BFS), and Depth-First Search (DFS), under increasing graph sizes, using adjacency matrices.

Due to the memory-intensive nature of storing full adjacency matrices, a hard cap emerged around 4000 nodes. Beyond this point, the Raspberry Pis lacked sufficient memory and killed the process. Therefore, this is the experiment for which there is least amount of accuracy when considering the usage of Pis as an analogue for real HPC nodes. Energy usage remained low across all three algorithms, even as the number of nodes scaled. This reinforces the idea that well-designed algorithms with efficient data access patterns can perform admirably to give energy savings.

502
503
504
505
506
507
508
509
510
511
512
513

Table 15. Dijkstra Performance and Energy Metrics

N	Time (s)	Energy (Wh)
10	1.1789	0.000128
20	1.2052	0.000160
30	1.1849	0.000128
40	1.2007	0.000128
50	1.1954	0.000160
60	1.1992	0.000160
70	1.1834	0.000160
80	1.1999	0.000128
90	1.1995	0.000096
100	1.1959	0.000160
150	1.1965	0.000128
200	1.2290	0.000160
250	1.2322	0.000128
500	1.4479	0.000160
750	1.6996	0.000287
1000	2.0021	0.000256
1500	2.9090	0.000447
2000	4.2382	0.000575
2500	5.9679	0.000831
3000	8.0646	0.001086
3500	12.5270	0.001661
4000	Killed	—


Table 16. BFS Performance and Energy Metrics

N	Time (s)	Energy (Wh)
10	2.5136	0.000383
20	2.1134	0.000224
30	1.1911	0.000128
40	1.2128	0.000128
50	1.1862	0.000096
60	1.2095	0.000096
70	1.1935	0.000096
80	1.1874	0.000128
90	1.1957	0.000096
100	1.2112	0.000096
150	1.1629	0.000096
200	1.2360	0.000128
250	1.2929	0.000096
500	1.3342	0.000096
750	1.5888	0.000224
1000	1.9163	0.000256
1500	2.7859	0.000383
2000	3.9720	0.000543
2500	5.9881	0.000831
3000	7.6786	0.001054
3500	11.9069	0.001565
4000	Killed	—

Table 17. DFS Performance and Energy Metrics

N	Time (s)	Energy (Wh)
10	1.1534	0.000128
20	1.1674	0.000096
30	1.1405	0.000128
40	1.1411	0.000096
50	1.1619	0.000128
60	1.1507	0.000096
70	1.1452	0.000096
80	1.1554	0.000128
90	1.1531	0.000128
100	1.2112	0.000096
150	1.1565	0.000096
200	1.1732	0.000128
250	1.2010	0.000128
500	1.3737	0.000128
750	1.5797	0.000256
1000	2.0100	0.000287
1500	2.7529	0.000383
2000	3.9876	0.000511
2500	5.5474	0.000767
3000	7.4741	0.000990
3500	9.7381	0.001438
4000	Killed	–

Execution Time of Graph Algorithms vs Number of Nodes

**Figure 17.** Execution Time vs Number of Nodes for Dijkstra, BFS, and DFS




Figure 18. Energy Usage of Dijkstra, BFS, and DFS with Increasing Graph Size

4.2. Analysis

This research confirms that energy efficiency and performance are deeply linked but not always in straightforward ways. Across all experiments, performance optimisations, whether via compiled languages, parallelisation, or overclocking, generally led to improved energy efficiency. Notably, the Monte Carlo Pi workload achieved a 4 \times improvement in EPS when moving from Python to C, and over 5 \times EPS when parallelised across 12 cores. These results underscore the impact of well-matched workloads and system configurations.

However, the study also highlights that speed and efficiency are not always synonymous. For workloads like recursive Fibonacci, increased parallelism led to worse energy metrics, showing that not all workloads when parallelised lead to gains which justify their energy cost. This pattern repeated with clocking experiments: overclocking offered significant speedups and energy gains up to a point, beyond which system instability prevented further gains.

Importantly, system-level strategies like cooling and on-demand computing demonstrated meaningful improvements, especially under heavy workloads. Cooling mitigated thermal throttling, resulting in smoother scaling, while on-demand computing significantly reduced overall energy usage.

Crucially, this study reveals an under-appreciated cost in energy-aware software development: developer time and skill. Writing parallel, efficient code for workloads that are not embarrassingly parallelisable is difficult. The benefit of such optimisation is only worthwhile when the workload is performed frequently or at scale, otherwise, the human and financial cost of optimisation may outweigh the energy savings. While our work focusses on small clusters and is not completely generalisable the same challenge exists for real-world HPC and green computing: balancing development overhead with sustainability goals.

Moreover, the findings in this project emphasise the need to embed energy awareness into the early stages of system design and software development. As energy costs rise and environmental concerns grow, future computing systems, from data centres to edge devices, will need to be optimised not only for performance, but for energy efficiency at every layer. This research reinforces the value of empirical testing over assumptions, showing that even small-scale, low-power clusters can reveal meaningful insights about workload behaviour. Encouragingly, it also shows that many gains can be achieved through

software and configuration alone, offering a low-barrier entry point to greener computing without requiring expensive new hardware.

5. Conclusions

5.1. Summary of Findings

- **Programming Language:** The choice of programming language had a notable impact on energy efficiency and execution time. C, a compiled language, significantly outperformed Python in both metrics, demonstrating that lower-level languages offer measurable energy savings, particularly for repetitive or computationally intensive workloads like Monte Carlo Pi.
- **Parallelisation and Multi-node Execution:** Parallelisation greatly improved both performance and energy efficiency—but only when applied to suitable workloads. Monte Carlo Pi, being embarrassingly parallel, showed strong scaling across up to 12 cores, with increased EPS and EF. However, recursive Fibonacci, a sequential workload, suffered from increased energy usage and reduced EPS when parallelised, showcasing the importance of workload profiling before parallel implementation.
- **Cooling:** Cooling hardware led to noticeable improvements in both runtime and energy usage under high loads. For large workloads ($N = 10^9$ and $N = 10^{10}$), cooled systems exhibited better EPS and EF compared to uncooled ones, confirming that thermal throttling on small devices like Raspberry Pis can significantly impact performance and efficiency.
- **Clock Speed:** Overclocking provided clear benefits up to a safe threshold (1500 MHz), improving execution time while maintaining or even improving energy efficiency. Higher speeds reduced the time processors were active, which outweighed the increase in instantaneous power draw—making overclocking an effective, low-effort optimisation.
- **EF & EPS:** EF and EPS proved to be reliable metrics for comparing energy-performance trade-offs. Across most experiments, particularly for Monte Carlo and sorting workloads, EF aligned well with speedup trends. However, in on-demand configurations or parallelising inefficient workloads, EF and EPS exposed energy waste hidden by performance metrics alone.

5.2. Key Takeaways

- Energy efficiency and performance are not always aligned. Effective optimisation requires matching system design to workload characteristics.
- Parallelisation offers strong gains for suitable workloads, but wastes energy if applied blindly.
- Simple system-level interventions, like cooling and clock tuning, offer tangible energy savings with little development overhead.
- Programming languages and algorithmic choices impact energy usage as much as runtime, emphasising the importance of considering both hardware and software layers in sustainable system design.
- Metrics like EF and EPS are essential in comparing energy-performance trade-offs and can reveal inefficiencies that raw speedup fails to capture.

5.3. Limitations and Potential Future Work

- This research was conducted over a limited timescale and future work could yield further insights through more detailed and extensive experimentation.
- The testbed consisted of three Raspberry Pi 3B+ nodes, which are not fully representative of the complexity, scale, and thermal behaviour of real-world HPC clusters or data centres.
- The energy data was interpreted with some support from real-world cost and emission statistics which, while taken from reputable sources, could not all be explicitly verified or cited.

- In the future, these experiments should be replicated on a larger HPC cluster to validate trends and scaling behaviours more accurately. 597
- Exploration of heterogeneous architectures—such as GPUs, FPGAs, and emerging technologies like quantum accelerators—would offer valuable insight into whether they can complete specific workloads in a more energy-efficient manner than traditional CPUs. 598
- Future work could also investigate automated tools for software-level optimisation, including compilers or profiling tools that recommend low-energy implementations. 599

Acknowledgements

The authors acknowledge the open-source community and the teams at DFRobot and Raspberry Pi, whose tutorials and dedication to accessible technology greatly supported this work.

1. International Energy Agency. Data Centres and Data Transmission Networks - Tracking Report, 2024. 600
2. van Kempen, N.; Kwon, H.J.; Nguyen, D.T.; Berger, E.D. It's Not Easy Being Green: On the Energy Efficiency of Programming Languages, 2025, [[arXiv:cs.PL/2410.05460](https://arxiv.org/abs/cs/PL/2410.05460)]. 601
3. Jiang, H.; Xiong, F.; Huang, Y. Energy efficiency improvement scheme based on edge computing. In Proceedings of the 2021 IEEE International Conference on Artificial Intelligence and Computer Applications (ICAICA), 2021, pp. 1033–1036. <https://doi.org/10.1109/ICAICA52286.2021.9498002>. 602
4. Masciari, E.; Napolitano, E.V. The Environmental Cost of High Performance Computing System Simulation. In Proceedings of the 2024 32nd Euromicro International Conference on Parallel, Distributed and Network-Based Processing (PDP), 2024, pp. 289–292. <https://doi.org/10.1109/PDP62718.2024.00048>. 603
5. Paul, S.G.; Saha, A.; Arefin, M.S.; Bhuiyan, T.; Biswas, A.A.; Reza, A.W.; Alotaibi, N.M.; Alyami, S.A.; Moni, M.A. A Comprehensive Review of Green Computing: Past, Present, and Future Research. *IEEE Access* **2023**, *11*, 87445–87494. <https://doi.org/10.1109/ACCESS.2023.3304332>. 604
6. Strohmaier, E.; Dongarra, J.; Simon, H.; Meuer, M. GREEN500 November 2024, 2024. 605
7. Department for Energy Security and Net Zero. Valuation of Energy Use and Greenhouse Gas Emissions. Technical report, UK Government, 2023. 606
8. International Energy Agency (IEA). CO₂ Emissions in 2023: Executive Summary, 2023. 607
9. Strubell, E.; Ganesh, A.; McCallum, A. Energy and policy considerations for modern deep learning research. In Proceedings of the Proceedings of the AAAI conference on artificial intelligence, 2020, Vol. 34, pp. 13693–13696. 608
10. Amazon Web Services. Carbon-free energy. <https://sustainability.aboutamazon.com/products-services/aws-cloud#:~:text=a%20new%20tab-,Carbon%2Dfree%20energy,of%20our%20original%202030%20goal>, 2025. Accessed: 2025-03-23. 609
11. Masciari, E.; Napolitano, E.V. Environmental Sustainability of AI: Estimating \$\$CO_2\$\$ Emissions Across Cloud, Edge, and Fog Paradigms. In Proceedings of the Web Information Systems Engineering – WISE 2024 PhD Symposium, Demos and Workshops; Barhamgi, M.; Wang, H.; Wang, X.; Aimeur, E.; Mrissa, M.; Chikhaoui, B.; Boukadi, K.; Grati, R.; Maamar, Z., Eds., Singapore, 2025; pp. 409–418. 610
12. Argerich, M.F.; Patiño-Martínez, M. Measuring and Improving the Energy Efficiency of Large Language Models Inference. *IEEE Access* **2024**, *12*, 80194–80207. <https://doi.org/10.1109/ACCESS.2024.3409745>. 611
13. Samsi, S.; Zhao, D.; McDonald, J.; Li, B.; Michaleas, A.; Jones, M.; Bergeron, W.; Kepner, J.; Tiwari, D.; Gadepally, V. From Words to Watts: Benchmarking the Energy Costs of Large Language Model Inference. In Proceedings of the 2023 IEEE High Performance Extreme Computing Conference (HPEC), 2023, pp. 1–9. <https://doi.org/10.1109/HPEC58863.2023.10363447>. 612
14. Eurostat. Electricity price statistics - Statistics Explained. https://ec.europa.eu/eurostat/statistics-explained/index.php?title=Electricity_price_statistics#Electricity_prices_for_non-household_consumers, 2024. 613
15. U.S. Bureau of Labor Statistics. Average Energy Prices, Selected Areas, Midwest Region. https://www.bls.gov/regions/midwest/data/averageenergyprices_selectedareas_table.htm, 2024. 614

16. Statista. China: Business electricity price 2024. <https://www.statista.com/statistics/1373596/business-electricity-price-china/>, 2024. 652
653
17. Rountree, B.; Ahn, D.H.; de Supinski, B.R.; Lowenthal, D.K.; Schulz, M. Beyond DVFS: A First Look at Performance under a Hardware-Enforced Power Bound. In Proceedings of the 2012 IEEE 26th International Parallel and Distributed Processing Symposium Workshops and PhD Forum, 2012, pp. 947–953. <https://doi.org/10.1109/IPDPSW.2012.116>. 654
655
656
657
18. Labasan, S. Energy-efficient and power-constrained techniques for exascale computing. *Semanticscholar: Seattle, WA, USA 2016*. 658
659
19. Wu, Y.; Mota, J.F.C.; Wallace, A.M. Joint Undervolting and Overclocking Power Scaling Approximation on FPGAs. In Proceedings of the 2022 Sensor Signal Processing for Defence Conference (SSPD), 2022, pp. 1–5. <https://doi.org/10.1109/SSPD54131.2022.9896229>. 660
661
662
20. Calore, E.; Gabbana, A.; Schifano, S.F.; Tripiccione, R. Evaluation of DVFS techniques on modern HPC processors and accelerators for energy-aware applications. *Concurrency and Computation: Practice and Experience* **2017**, 29. <https://doi.org/10.1002/cpe.4143>. 663
664
665
21. Ricciardi, S.; Santos-Boada, G.; Careglio, D.; Palmieri, F.; Fiore, U. Evaluating energy savings in WoL-enabled networks of PCs. In Proceedings of the 2013 IEEE International Symposium on Industrial Electronics, 2013, pp. 1–6. <https://doi.org/10.1109/ISIE.2013.6739357>. 666
667
668
22. Wang, J.; Feng, L.; Xue, W. A review of energy efficiency technology in computer servers and cluster systems. In Proceedings of the 2011 3rd International Conference on Computer Research and Development, 2011, Vol. 2, pp. 109–113. <https://doi.org/10.1109/ICCRD.2011.5764094>. 669
670
23. Qasaimeh, M.; Denolf, K.; Lo, J.; Vissers, K.; Zambreno, J.; Jones, P.H. Comparing energy efficiency of CPU, GPU and FPGA implementations for vision kernels. In Proceedings of the 2019 IEEE international conference on embedded software and systems (ICESS). IEEE, 2019, pp. 1–8. 672
673
674
675
24. Ma, K.; Li, X.; Chen, W.; Zhang, C.; Wang, X. Greengpu: A holistic approach to energy efficiency in gpu-cpu heterogeneous architectures. In Proceedings of the 2012 41st international conference on parallel processing. IEEE, 2012, pp. 48–57. 676
677
25. Jin, C.; de Supinski, B.R.; Abramson, D.; Poxon, H.; DeRose, L.; Dinh, M.N.; Endrei, M.; Jessup, E.R. A survey on software methods to improve the energy efficiency of parallel computing. *Int. J. High Perform. Comput. Appl.* **2017**, 31, 517–549. <https://doi.org/10.1177/1094342016665471>. 679
680
681
26. Sampson, A.; Dietl, W.; Fortuna, E.; Gnanapragasam, D.; Ceze, L.; Grossman, D. EnerJ: approximate data types for safe and general low-power computation. In Proceedings of the Proceedings of the 32nd ACM SIGPLAN Conference on Programming Language Design and Implementation, New York, NY, USA, 2011; PLDI ’11, p. 164–174. <https://doi.org/10.1145/1993498.1993518>. 682
683
684
685
686
27. Liu, C.; Qin, X.; Kulkarni, S.; Wang, C.; Li, S.; Manzanares, A.; Baskiyar, S. Distributed Energy-Efficient Scheduling for Data-Intensive Applications with Deadline Constraints on Data Grids. In Proceedings of the 2008 IEEE International Performance, Computing and Communications Conference, 2008, pp. 26–33. <https://doi.org/10.1109/PCCC.2008.4745123>. 687
688
689
690